

COMS 4772 Fall 2015: Homework #1

Daniel M Sheehan - dms2203@columbia.edu

Discussants: tudor

February 10, 2016

Problem 1

```
from scipy.io import loadmat
ocr = loadmat('ocr.mat')

import matplotlib.pyplot as plt
from matplotlib import cm

import random
import numpy as np

import datetime
import csv

def dist(x, y):
    diff = x.astype(np.float32) - y.astype(np.float32)
    return diff.dot(diff)

def nnClassifier(trainingData, trainingLabels, testData, testLabels):
    total_count = len(testLabels)
    wrong_count = 0
    for testVector, testLabel in zip(testData, testLabels):
        # find the closest training vector
        minDist = float('inf')
        minLabel = None
        for trainingVector, trainingLabel in zip(trainingData, trainingLabels):
            d = dist(testVector, trainingVector)
            #print 'd:', d, 'minDist:', minDist
            if d < minDist:
```

```

        minDist = d
        minLabel = trainingLabel

    if minLabel[0] != testLabel[0]:
        wrong_count += 1

    print 'wrongcount:', wrong_count
    errorRate = float(wrong_count) / total_count
    print errorRate
    return errorRate

testData = ocr['testdata']
testLabels = ocr['testlabels']

sample_sizes = [1000,2000,4000,8000]

for i in range(10):
    dt = datetime.datetime.now().strftime("%Y-%m-%d %H:%M:%S")
    dt = dt.replace(' ', '-').replace(':', '-')
    print dt

    error_rates = []
    for n in sample_sizes:
        print n
        sel = random.sample(xrange(60000),n)
        trainingData = ocr['data'][sel]
        trainingLabels = ocr['labels'][sel]
        e = nnClassifier(trainingData, trainingLabels, testData, testLabels)
        error_rates.append(e)

    print error_rates

    with open("data/output"+dt+".csv", "wb") as f:
        writer = csv.writer(f)
        row = sample_sizes, error_rates
        writer.writerows(row)

import glob, os
import pandas as pd

theFiles = glob.glob('data/*.csv')
df_list = []

for i in theFiles:

```

```

df = pd.read_csv(i)
df = df.T
df.columns = ['error_rate']
df['date'] = i.replace('data/output', '').replace('.csv', '')
df['samp'] = df.index
df_list.append(df)

```

```

df = pd.concat(df_list)
df.head(5)

```

	error_rate	date	samp
1000	0.1144	2016-02-05-11-00-00	1000
2000	0.0894	2016-02-05-11-00-00	2000
4000	0.0702	2016-02-05-11-00-00	4000
8000	0.0534	2016-02-05-11-00-00	8000
1000	0.1119	2016-02-05-11-41-39	1000

```

df['count'] = 1
dfg = df.groupby(['samp'],as_index=False).sum()
dfg['error_rate'] = dfg['error_rate'] / dfg['count']
dfg.head(10)

```

	samp	error_rate	count
0	1000	0.11471	10
1	2000	0.08724	10
2	4000	0.06922	10
3	8000	0.05479	10

```

dfgby = df.groupby(['samp']).std()
dfgby.head(10)

```

	error_rate	count
samp		
1000	0.003527	0
2000	0.003654	0
4000	0.001981	0
8000	0.002001	0

```
import matplotlib.pyplot as plt
plt.style.use('ggplot')
%matplotlib inline

#!/usr/bin/env python
import numpy as np
import matplotlib.pyplot as plt


x = dfg['samp']
y = dfg['error_rate']

theError = dfgby['error_rate']

# First illustrate basic pyplot interface, using defaults where possible.
plt.figure()
plt.errorbar(x, y, yerr=theError)#, xerr=0.2)
plt.title("Error Rate for 1-nearest neighbor classifier with Euclidean distance")
plt.ylim(ymin=0, ymax=0.13)
plt.xlim(xmin=0, xmax=9000)

(0, 9000)
```

Error Rate for 1-nearest neighbor classifier with Euclidean distance



Problem 2

$$P(X | class) = \prod_{j=1}^d \mu_j^{x_j} (1 - \mu_j)^{1-x_j}$$

$$L = \log(\dots) \rightarrow \left(\frac{\delta L}{\mu_j} \right) = \delta$$

$$P(motorcycle) P(gas) \dots P(mouse) \leftarrow O$$

$$\delta/N \rightarrow \left(\frac{1+O}{2+N} \right)$$

$$= (MLE | Laplace)$$

$$P(Y | X)$$

$$prediction = \operatorname{argmax}_{forum} P(Y = forum | X)$$

$$\propto P(X | forum) P(forum)$$

$$\Rightarrow \prod_{j=1}^d \mu_j^{x_j} (1 - \mu_j)^{1-x_j} \# \text{ forum/total}$$

$$P(data | \mu) \rightarrow P(data | \mu_{forum_1}), P(data | \mu_{forum_2})$$

$$Likelihood = \prod P(\text{word appears}) P(\text{!word appears})$$

$$\operatorname{argmax}_y \left(\prod_{j=1}^d \mu_{y,j}^{x_j} (1 - \mu_{y,j})^{1-x_j} \right) \pi_y$$

$$Linear \text{ classification} : y = w_o + w_1 x_1 + w_2 x_2 + \dots$$

$$y = mx + b = w_o + \sum_{j=1}^d w_j x_j$$

$$\log[(\prod_{j=1}^d \mu_{y,j}^{x_j} (1 - \mu_{y,j})^{1-x_j}) \pi_y]$$

$$= \log \pi_y + \sum_{j=1}^d x_j \log \mu_{y,j} + (1 - x_j) \log(1 - \mu_{y,j})$$

$$= \log \pi_y + \sum_j \log(1 - \mu_{i,j}) + \sum_j [\log \mu_{j,y} * \log(1 - \mu_{i,j})] x_j$$

$$= b_y + w_y x_j \rightarrow w_{y(\text{forum or class}), j(\text{dimension or word index})}$$

$$\prod_{j=1} \mu \dots (1 - \mu) \prod_y$$

$$b + wx(\text{dotproduct})$$

```
from scipy.io import loadmat
import numpy as np
import pandas as pd

news = loadmat('news.mat')

def get_params(Xtr, Ytr):
    # finding w and b for each class
    Ytr = Ytr.flatten()
    theClasses = set(Ytr) #list of classes
    totCnt = len(Ytr)
    paramDict = {}
    for theClass in theClasses:
        idx = np.where(Ytr==theClass)[0]
        cnt = len(idx)
        xForThisClass = Xtr[idx]
        mu = (1 + xForThisClass.sum(axis=0)) / ( 2 + cnt)
        pi = cnt/float(totCnt)
        bias = np.log(pi) + np.log(1 - mu).sum() #lgo pi + sum of all log minus meu
        weights = np.log(mu) - np.log(1 - mu) #log meu, log 1 - meu
        paramDict[theClass] = {'bias':bias, 'weights':weights}
    return paramDict

def get_error_rate(params, Xtest, Ytest):
    total_count = len(Ytest)
    wrong_count = 0
    for testVector, testLabel in zip(Xtest, Ytest):
        #print testVector
        x = testVector
        maxLL = float('-inf')
```

```

minLabel = None
for key, value in params.iteritems():
    # make a prediction for input testVector
    w = value['weights']
    b = value['bias']
    # print "weights shape:", w.shape, "bias shape:", b.shape, "x shape:", x.shape
    # # np.dot(w, x) ##### no! x is not a numpy array
    # # w.dot(x) ##### x is not a numpy array
    # # np.dot(w, x.to_array()) ##### slow
    # # w.dot(x.to_array()) ##### same
    # # x.dot(w)
    pred = x.dot(w.T) + b #should pred label 1-20?
    # if this prediction != testLabel, increment wrong_count
    if pred > maxLL:
        maxLL = pred
        minLabel = key
if int(minLabel) != int(testLabel[0]):
    wrong_count += 1

```

```

#print 'wrongcount:', wrong_count
errorRate = float(wrong_count) / total_count
#print errorRate
return errorRate
#which class has the highest w*x + b

```

```

Xtrain = news['data']
Ytrain = news['labels']
params = get_params(Xtrain, Ytrain)
print "Training error:", get_error_rate(params, Xtrain, Ytrain)

```

```

Xtest = news['testdata']
Ytest = news['testlabels']
print "Test error:", get_error_rate(params, Xtest, Ytest)

```

```

#-----

```

```

dfv = pd.read_csv("news.vocab",header=None)
dfv.columns = ['word']
dfv['vuid'] = dfv.index + 1

```

```

dfg = pd.read_csv("news.groups",header=None)
dfg.columns = ['topic'] #just declare col names here
dfg['topic'] = dfg.topic.str.split(' ',1).str[0]
dfg['guid'] = dfg.index + 1

```

```

df = pd.DataFrame(params)
df = df.T
df['guid'] = df.index

#print df.head(20) #print dfv.head(20) #print dfg.head(20)

df = df.merge(dfg, on='guid', how='left')
#print df.head(20)

dfw = df[['weights']]
dfw = dfw['weights'].tolist()

df_list = []

for i, k in enumerate(dfw):
    j = k.tolist()
    df = pd.DataFrame(j)
    df = df.T
    df.columns = ['weights']
    df['guid'] = i + 1
    df['vuid'] = df.index + 1
    df = df.sort('weights', ascending=False).head(20)
    #print df.head(25)
    df_list.append(df)

df = pd.concat(df_list)
df = df.merge(dfv, on='vuid', how='left').merge(dfg, on='guid', how='left')
print df.head(400)

df.to_csv('hw1-2.csv', index=False)

```

Training error: 0.216256988198

Test error: 0.37601598934

	weights	guid	vuid	word	topic
0	3.854394	1	29	the	alt.atheism
1	2.516800	1	12	of	alt.atheism
2	2.487156	1	30	in	alt.atheism
3	2.375473	1	33	to	alt.atheism
4	2.070139	1	23	and	alt.atheism
5	1.989290	1	233	that	alt.atheism
6	1.755913	1	60	is	alt.atheism
7	1.457848	1	778	writes	alt.atheism
8	1.404548	1	42	it	alt.atheism
9	1.266150	1	474	you	alt.atheism
10	0.903257	1	722	not	alt.atheism

11	0.883123	1	144	be	alt.atheism
12	0.775212	1	775	edu	alt.atheism
13	0.756061	1	81	for	alt.atheism
14	0.746535	1	251	this	alt.atheism
15	0.671441	1	27	are	alt.atheism
16	0.652956	1	922	have	alt.atheism
17	0.508613	1	51	but	alt.atheism
18	0.490957	1	48	on	alt.atheism
19	0.464613	1	473	if	alt.atheism
20	1.897778	2	29	the	comp.graphics
21	1.539057	2	33	to	comp.graphics
22	1.161351	2	23	and	comp.graphics
23	1.123930	2	30	in	comp.graphics
24	1.105485	2	12	of	comp.graphics
25	0.929536	2	81	for	comp.graphics
26	0.774026	2	60	is	comp.graphics
27	0.596870	2	42	it	comp.graphics
28	0.182322	2	233	that	comp.graphics
29	0.106447	2	48	on	comp.graphics
..
370	1.070201	19	474	you	talk.politics.misc
371	0.960271	19	251	this	talk.politics.misc
372	0.938933	19	722	not	talk.politics.misc
373	0.917793	19	27	are	talk.politics.misc
374	0.896844	19	770	article	talk.politics.misc
375	0.696368	19	922	have	talk.politics.misc
376	0.657944	19	144	be	talk.politics.misc
377	0.629431	19	48	on	talk.politics.misc
378	0.426971	19	52	with	talk.politics.misc
379	0.400103	19	388	as	talk.politics.misc
380	3.186353	20	29	the	talk.religion.misc
381	2.251292	20	12	of	talk.religion.misc
382	2.191359	20	30	in	talk.religion.misc
383	2.162438	20	33	to	talk.religion.misc
384	1.976494	20	60	is	talk.religion.misc
385	1.904237	20	23	and	talk.religion.misc
386	1.517065	20	233	that	talk.religion.misc
387	1.008897	20	42	it	talk.religion.misc
388	0.942363	20	722	not	talk.religion.misc
389	0.765468	20	474	you	talk.religion.misc
390	0.657640	20	778	writes	talk.religion.misc
391	0.622530	20	922	have	talk.religion.misc
392	0.610909	20	81	for	talk.religion.misc
393	0.599328	20	27	are	talk.religion.misc
394	0.530628	20	251	this	talk.religion.misc

```

395  0.519300    20   144      be  talk.religion.misc
396  0.463131    20    48      on  talk.religion.misc
397  0.363747    20    52    with talk.religion.misc
398  0.352821    20    51      but talk.religion.misc
399  0.331033    20   770  article talk.religion.misc

```

```
[400 rows x 5 columns]
```

```
/usr/local/lib/python2.7/site-packages/IPython/kernel/__main__.py:98: FutureWarning: sort(columns=)
```

Problem 3

$$P(Yes_{Y=1}) = 0.001$$

$$\pi_1 = "$$

$$P(Yes_{Y=0}) = 0.999$$

$$\pi_0 = "$$

$$\pi_1 = 200/300 = 2/3 \quad \pi_2 = 100/300 = 1/3$$

$$not \rightarrow P(X|C_1) \text{ vs. } P(X|C_2)$$

$$V \rightarrow P(C_1|X) \text{ vs. } P(C_2|X)$$

$$P(X|C_1)\pi_1 = P(X|C_2)\pi_2 \quad C$$

$$* 2/3 \quad * 1/3$$

$$P(X|C_1)\pi_1 = P(X|C_2)\pi_2 \quad C$$

$$\text{find } x$$

$$\pi_0 = 2/3 \quad \pi_1 = 1/3$$

$$P(X|Y = 0) = N(0, 1) = \frac{1}{\sqrt{2\pi}} e^{-1/2 * x^2}$$

$$P(X|Y = 1) = N(1, 1/4) = \frac{1}{\sqrt{2\pi * 1/4}} e^{-1/2 * (x-1)^2 / (1/4)}$$

$$P(X|Y = 0)\pi_0 = P(X|Y = 1)\pi_1 C$$

$$f * (x) = \begin{cases} 0 & x < b_1 \\ 1 & b_1 < x < b_2 \\ 0 & x > b_2 \end{cases}$$

Problem 4

a

$$C = \{red, orange, yellow, green, blue\}$$

$$n = 100 \text{ balls}$$

$$n_{red}, n_{orange}, n_{yellow}, n_{green}, n_{blue}$$

$$p(red) = n_{red}/n$$

$$p(red, red) = n_{red}/n * n_{red}/n$$

$$p(differentcolor) = 1 - p(samecolor)$$

$$= 1 - [p(red, red) + p(orange, orange) + p(yellow, yellow) + p(green, green) + p(blue, blue)]$$

$$= 1 - \sum_{c \in C} (n_c/n)^2$$

$$= 1 - \sum_c (p_c)^2$$

b

$$entropy H = -\{p \log p + (1 - p) \log(1 - p)\}$$

$$\max H \text{ wrt } p \rightarrow p = 0.5$$

$$H = - \sum_k P_k \log p_k$$

Paint each color 20 times.

$$\sigma(\omega_0 + \omega_1 x_1 + \omega_2 x_2 + \dots) \rightarrow [0, 1]$$

$$P(Y_{class} = 1 | X_{data})$$

posterior

$$P(Y = 1|X) = \frac{P(X, Y = 1)}{P(X)}$$

$$= \frac{P(X|Y = 1)P(Y)}{P(X)}$$

$$\propto P(X|Y = 1)P(Y)$$

likelihood – prior

$$P(Yes_{Y=1}) = 0.001$$

$$\pi_1 = ''$$

$$P(Yes_{Y=0}) = 0.999$$

$$\pi_0 = ''$$

$$\pi_1 = 200/300 = 2/3 \quad \pi_2 = 100/300 = 1/3$$

$$not \rightarrow P(X|C_1) \text{ vs. } P(X|C_2)$$

$$V \rightarrow P(C_1|X) \text{ vs. } P(C_2|X)$$

$$P(X|C_1)\pi_1 = P(X|C_2)\pi_2 \quad C$$

$$* 2/3 \qquad * 1/3$$

$$P(X|C_1)\pi_1 = P(X|C_2)\pi_2 \quad C$$

find x

$$\pi_0 = 2/3 \quad \pi_1 = 1/3$$

$$P(X|Y = 0) = N(0, 1) = \frac{1}{\sqrt{2\pi}} e^{-1/2 * x^2}$$

$$P(X|Y = 1) = N(1, 1/4) = \frac{1}{\sqrt{2\pi * 1/4}} e^{-1/2 * (x-1)^2 / (1/4)}$$

$$P(X|Y = 0)\pi_0 = P(X|Y = 1)\pi_1 \quad C$$

$$f \ast (x) = \begin{cases} 0 & x < b_1 \\ 1 & b_1 < x < b_2 \\ 0 & x > b_2 \end{cases}$$